

09/10/2020
LAB:11

Experiments using variational Autoencoder

* Aim =

To implement a variational autoencoder for compressing and generating MNIST digit images using a probabilistic approach.

* Objective =

1. To learn a latent space representation of MNIST digits.
2. To generate new images by sampling from the learned latent distribution.
3. To compare the reconstructed and generated images with the original dataset.

* Pseudocode =

Begin

Load MNIST dataset

Normalize and flatten images ($28 \times 28 \rightarrow 784$)

Define encoder:

Input layer \rightarrow Dense (256, ReLU)

mean layer(μ) \rightarrow Dense (latent-dim)

log variance layer(σ^2) \rightarrow Dense (latent-dim)

Sample latent vector $z = \mu + \epsilon \exp(\sigma^2/2) * \epsilon$ where $\epsilon \sim \mathcal{N}(0, 1)$

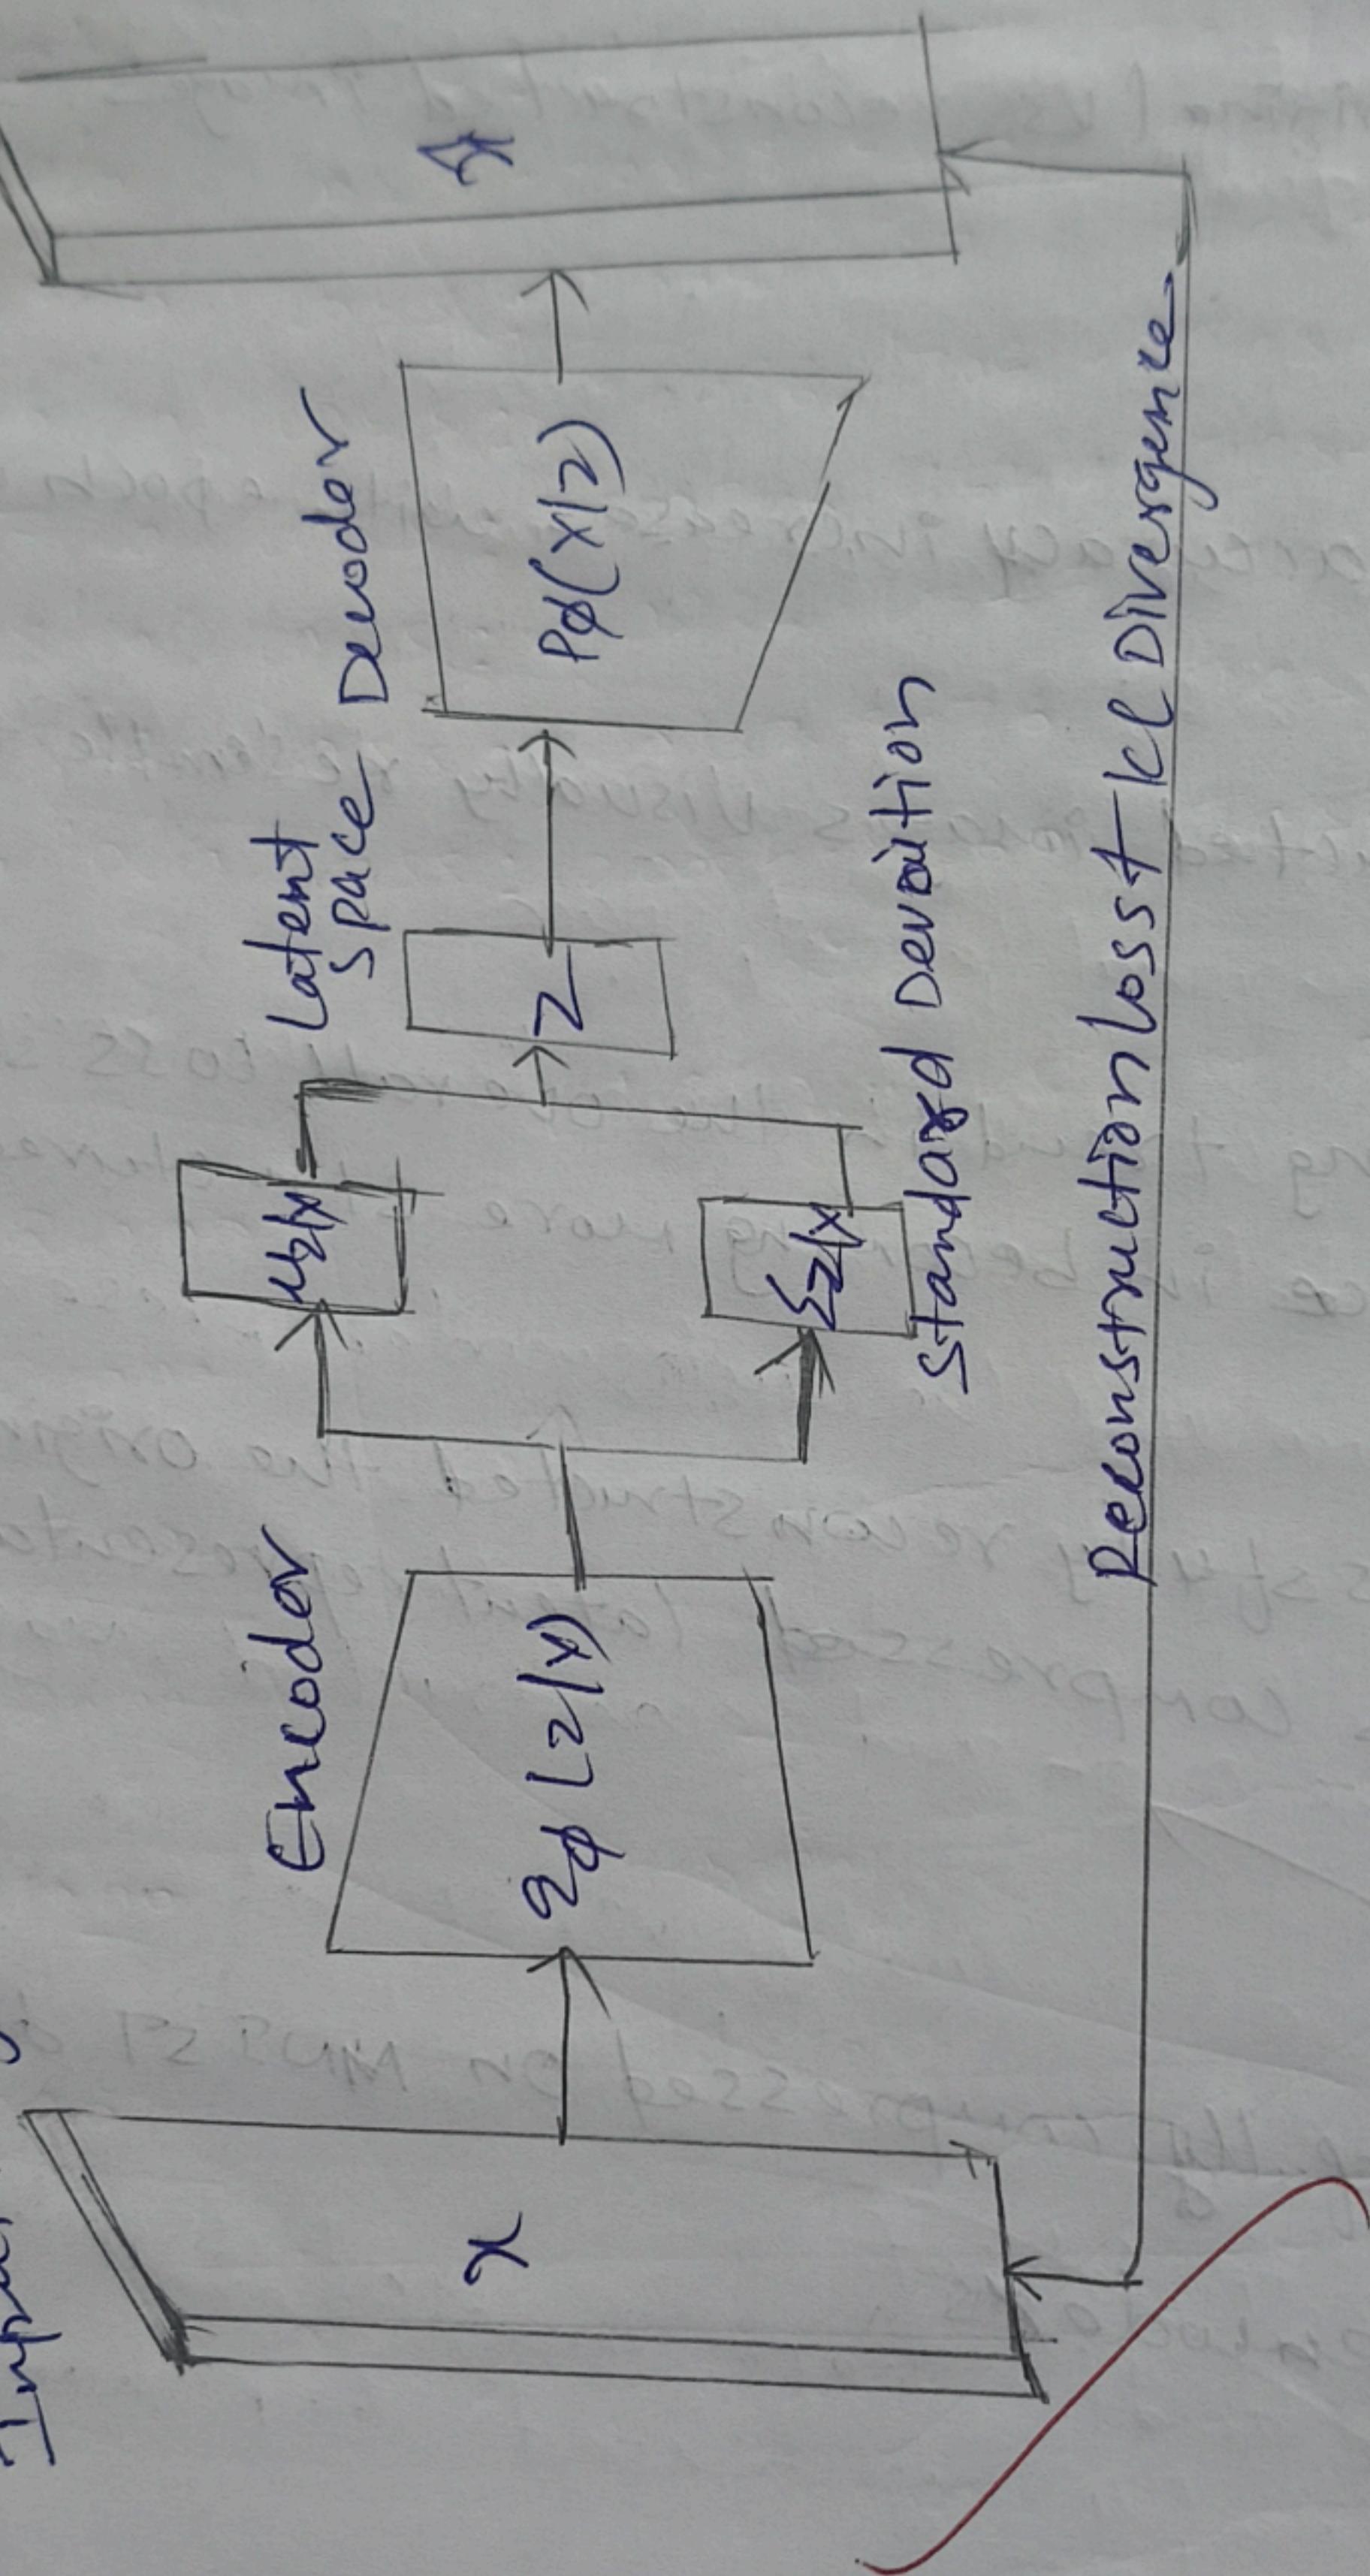
Define decoder:

Input $z \rightarrow$ Dense (256, ReLU)

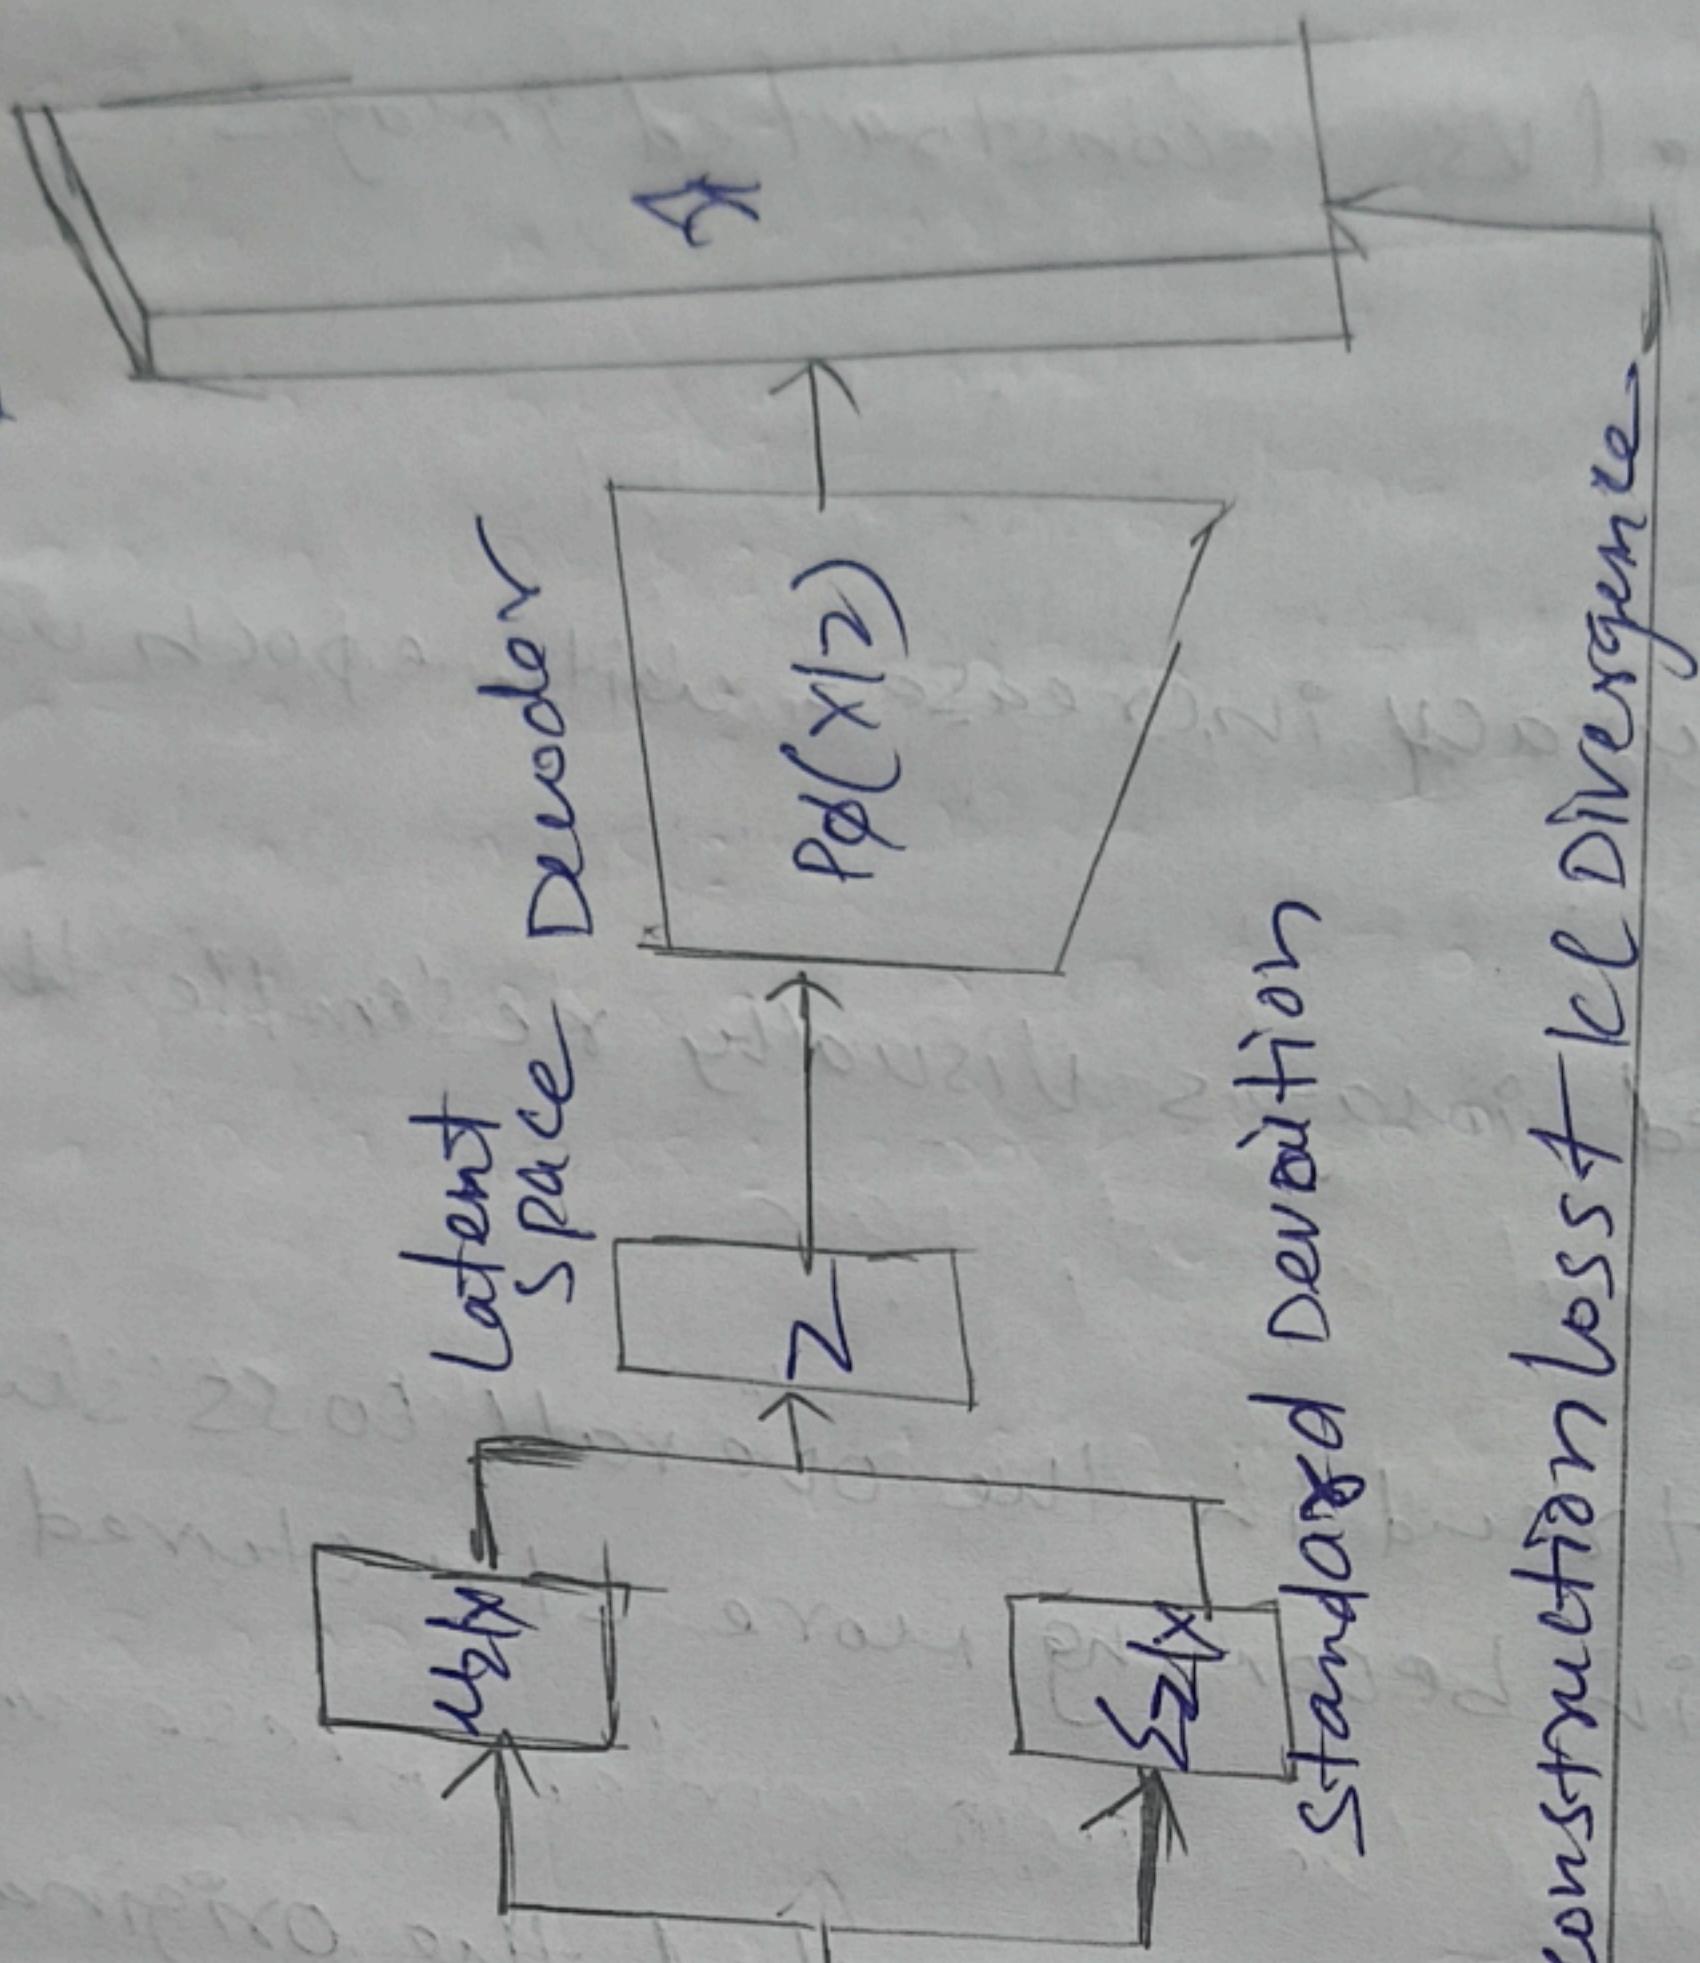
Output \rightarrow Dense (784, Sigmoid) \rightarrow reconstructed image

Architecture of VAT

Input image



Reconstructed image

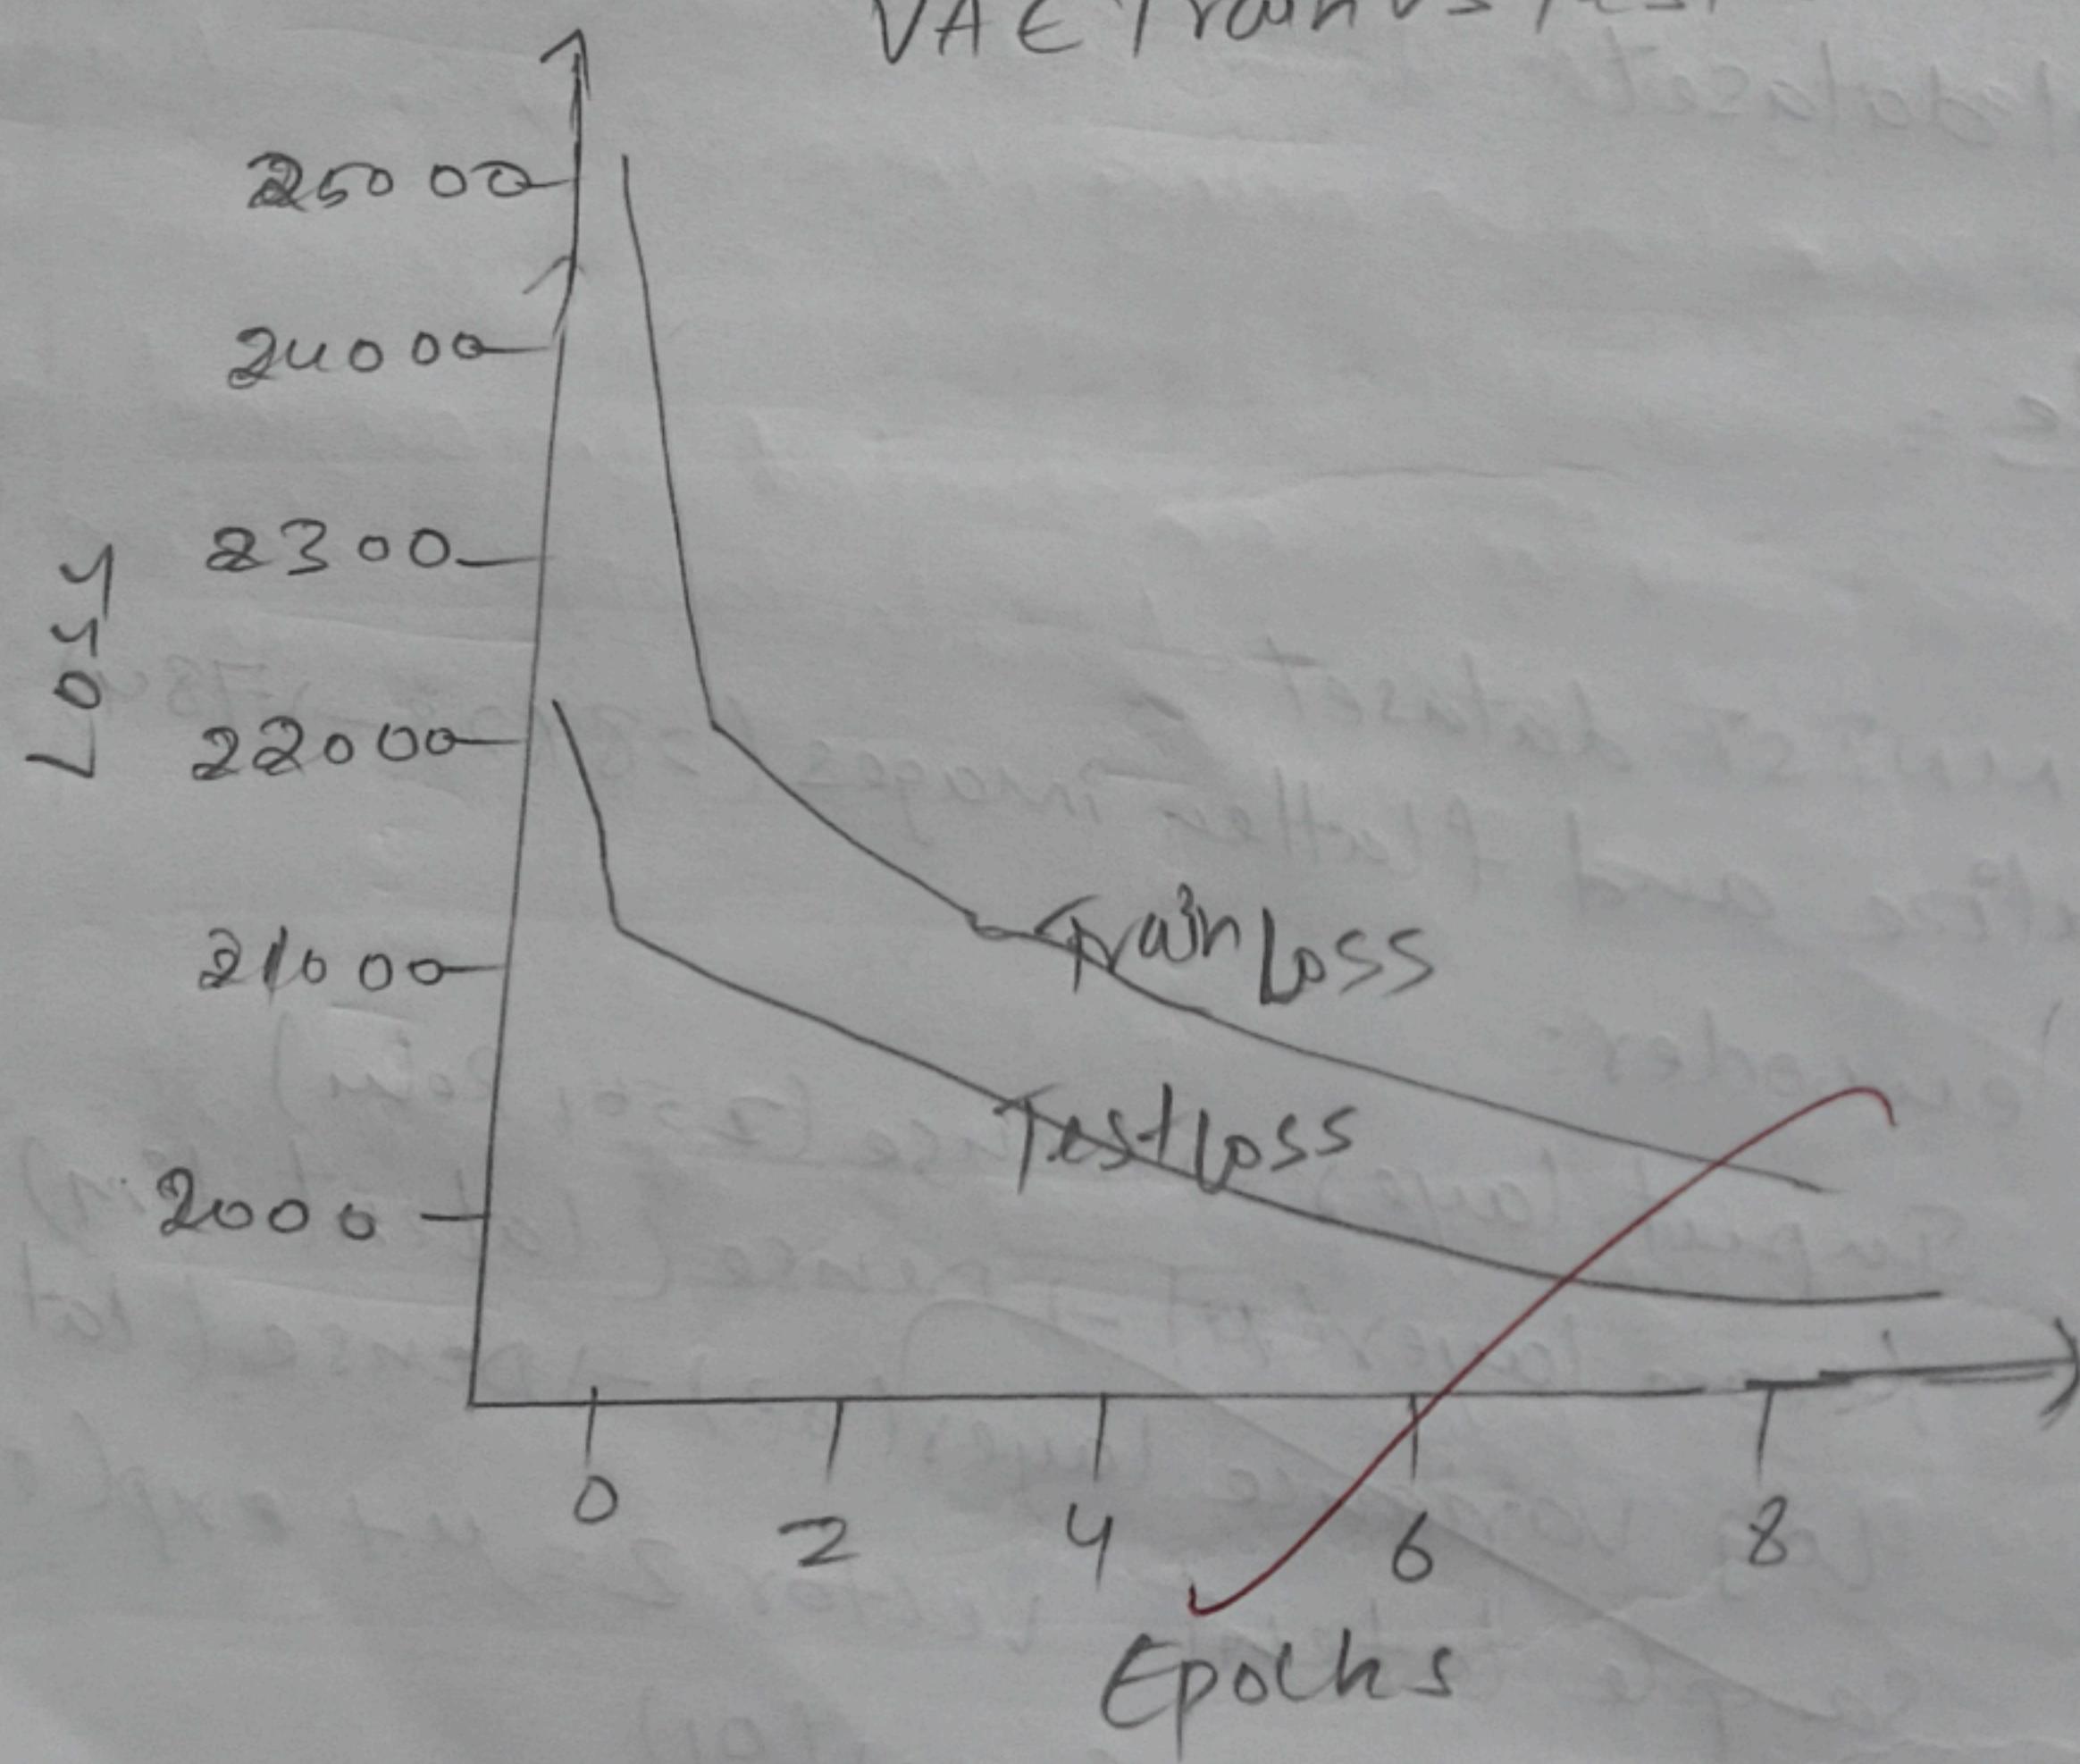


Reconstruction loss & KL Divergence

Output:-

- Epoch 1/10, loss: 25003.6851, Test loss: 22058.4611
Epoch 2/10, loss: 21878.4032, Test loss: 21122.0819
Epoch 3/10, loss: 21232.1173, Test loss: 20715.5481
Epoch 4/10, loss: 20888.5306, Test loss: 20482.5161
Epoch 5/10, loss: 20659.487, Test loss: 20302.6661
Epoch 6/10, loss: 20486.1709, Test loss: 20039.9918
Epoch 7/10, loss: 20346.4544, Test loss: 19970.3017
Epoch 8/10, loss: 20235.4947, Test loss: 19922.5311
Epoch 9/10, loss: 20137.8976, Test loss: 19820.2514
Epoch 10/10, loss: 20055.5366, Test loss: 19820.2514

VAE Train vs Test loss



Define loss:
Total loss = Reconstruction loss + KL divergence loss
Train model using Adam optimizer for multiple epochs
Encode test images \rightarrow get latent vectors
Decode latent vectors \rightarrow reconstruct images
Sample random z values \rightarrow generate new images

end

*Observation =

- Instead of mapping input to a fixed point, VAE encode input into a probability distribution in latent space
- VAE's use the reparameterization trick to allow for back propagation through the sampling process
- The loss function of VAE includes a reconstruction loss & KL divergence to regularize space & encourage to follow a prior ~~distribution~~

*Result = Successfully implemented VAE (Variational autoencoder) using MNIST dataset.

B.G.

```
[ ] ⏎ import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import numpy as np

# -----
# Hyperparameters
# -----
batch_size = 128
epochs = 10
learning_rate = 1e-3
latent_dim = 2 # 2D latent space for visualization

# -----
# Dataset
# -----
transform = transforms.ToTensor()
train_dataset = datasets.MNIST(root='data', train=True, download=True, transform=transform)
test_dataset = datasets.MNIST(root='data', train=False, download=True, transform=transform)

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)

# -----
# VAE Model
# -----
class VAE(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(28*28, 256)
        self.fc_mu = nn.Linear(256, latent_dim)
        self.fc_logvar = nn.Linear(256, latent_dim)
        self.fc2 = nn.Linear(latent_dim, 256)
        self.fc3 = nn.Linear(256, 28*28)

    def encode(self, x):
        h = torch.relu(self.fc1(x))
        return self.fc_mu(h), self.fc_logvar(h)

    def reparameterize(self, mu, logvar):
        std = torch.exp(0.5*logvar)
        eps = torch.randn_like(std)
        return mu + eps * std
```



```
[ ] ⏎ def forward(self, x):
    mu, logvar = self.encode(x)
    z = self.reparameterize(mu, logvar)
    return self.decode(z), mu, logvar

# -----
# Device and Model
# -----
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = VAE().to(device)
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

# -----
# Loss Function
# -----
def loss_fn(recon_x, x, mu, logvar):
    BCE = nn.functional.binary_cross_entropy(recon_x, x.view(-1, 28*28), reduction='sum')
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    return BCE + KLD

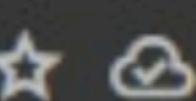
# -----
# Training Loop with Train & Test Loss
# -----
train_losses = []
test_losses = []

for epoch in range(epochs):
    model.train()
    running_loss = 0
    for x, _ in train_loader:
        x = x.to(device).view(x.size(0), -1) # Flatten
        optimizer.zero_grad()
        recon, mu, logvar = model(x)
        loss = loss_fn(recon, x, mu, logvar)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
    avg_train_loss = running_loss / len(train_loader)
    train_losses.append(avg_train_loss)

    # Compute test loss
    model.eval()
    test_loss = 0
    with torch.no_grad():
        for x, _ in test_loader:
            x = x.to(device).view(x.size(0), -1)
            recon, mu, logvar = model(x)
            test_loss += loss_fn(recon, x, mu, logvar).item()
    avg_test_loss = test_loss / len(test_loader)
```



Lab10.ipynb



File Edit View Insert Runtime Tools Help

Commands

```
[ ] # Visualize 2D Latent Space  
# -----  
x_sample, y_sample = next(iter(test_loader))  
x_sample = x_sample.to(device).view(x_sample.size(0), -1)  
with torch.no_grad():  
    mu, logvar = model.encode(x_sample)  
mu = mu.cpu().numpy()  
  
plt.figure(figsize=(8,6))  
scatter = plt.scatter(mu[:,0], mu[:,1], c=y_sample, cmap=cm.tab10)  
plt.colorbar(scatter)  
plt.title('VAE 2D Latent Space (Test Set)')  
plt.show()
```

```
Epoch 1/10 | Train Loss: 25003.6881 | Test Loss: 22056.4637  
Epoch 2/10 | Train Loss: 21878.4032 | Test Loss: 21122.0819  
Epoch 3/10 | Train Loss: 21232.1173 | Test Loss: 20715.5480  
Epoch 4/10 | Train Loss: 20888.5306 | Test Loss: 20482.5152  
Epoch 5/10 | Train Loss: 20658.4827 | Test Loss: 20302.1048  
Epoch 6/10 | Train Loss: 20486.1709 | Test Loss: 20138.8614  
Epoch 7/10 | Train Loss: 20346.4544 | Test Loss: 20039.9918  
Epoch 8/10 | Train Loss: 20235.4947 | Test Loss: 19970.3085  
Epoch 9/10 | Train Loss: 20137.8976 | Test Loss: 19922.5389  
Epoch 10/10 | Train Loss: 20055.5366 | Test Loss: 19820.2528
```

VAE Train vs Test Loss

