

~~Confidential~~

D. Khoja Nawaz

RA2311047060037

Deep Learning Techniques (Lab)

| Date       | Title  | Sign             |
|------------|--|------------------|
| 24/07/2025 | 1. Exploring the deep learning platform  | <del>Att'd</del> |
| 31/07/2025 | 2. Implement a classifier using open-source dataset                              | <del>Att'd</del> |
| 31/07/2025 | 3. Study of the classifiers with respect to statistical parameters               | <del>Att'd</del> |
| 14/08/2025 | 4. Build a simple feed forward neural network to recognize handwritten character | <del>Att'd</del> |
| 22/08/2025 | 5. Study of Activation functions and their role                                  | <del>Att'd</del> |
| 09/09/2025 | 6. Implement gradient descent and backpropagation in deep neural network         | <del>Att'd</del> |
| 16/09/2025 | 7. Build a CNN model to classify cat and Dog image                               | <del>Att'd</del> |
| 30/09/2025 | 8. Experiment of LSTM  | <del>Att'd</del> |
| 30/09/2025 | 9. Build a Recurrent Neural Network  | <del>Att'd</del> |
| 9/10/2025  | 10. Perform compression on MNIST dataset using autoencoder                       | <del>Att'd</del> |
| 9/10/2025  | 11. Experiments using variational autoencoder                                    | <del>Att'd</del> |
| 03/11/2025 | 12. Implement a Deep convolutional GAN to generate complex coloring              | <del>Att'd</del> |
| 03/11/2025 | 13. Understanding the architecture of pre-trained model                          | <del>Att'd</del> |
| 03/11/2025 | 14. Implement a pre-trained CNN model as feature extractor using                 | <del>Att'd</del> |
| 03/11/2025 | 15. Implement a YOLO model to detect objects                                     | <del>Att'd</del> |

03/11/25 LAB:12

Implement a Deep Convolutional GAN to generate complex color image

\*AIM =

To implement a Deep convolutional Generative Adversarial Network (DCGAN) of generating complex RGB color image that resemble real-world images.

\*Objective =

1. To understand the working of Generative Adversarial network (GANs)
2. To use convolutional layers for image generation and discrimination.
3. To train the Generator to produce realistic image and the Discriminator to distinguish real from fake
4. To evaluate the quality of generated images visually after training.

\* pseudocode =

Begin

1. Load dataset of real color images  
→ Resize and normalize images to (-1, 1)

2 Define Generator network (G):

Input: random noise vector (z)

Layers: series of convTranspose<sup>2D + Batch norm + ReLU</sup>

Output: RGB image


3. Define discriminator network (D):

Input: RGB image

Layers: series of conv2D + Batch norm + Leaky ReLU

Output: single probability

## Architecture of GAN



output

Epoch (1/10), D loss : 0.0992, G loss : 0.5678

Epoch (2/10), D loss : 0.5860, G loss : 2.7113

Epoch (3/10), D loss : 1.2431, G loss : 3.9827

Epoch (4/10), D loss : 0.1285, G loss : 4.1865

Epoch (5/10), D loss : 0.1983, G loss : 3.7253


Epoch (6/10), D loss : 0.2754, G loss : 4.0325

Epoch (7/10), D loss : 0.1110, G loss : 8.6499

Epoch (8/10), D loss : 0.0016, G loss : 8.4825

Epoch (9/10), D loss : 0.1210, G loss : 4.8809

Epoch (10/10), D loss : 0.0389, G loss : 6.7671



4. SET loss function = Binary cross entropy  
SET optimizes = adam ( $\gamma = 0.0002$ ,  $B_1 = 0.5$ )

5. For each epoch:

a. Train Discriminator:

- Feed real image  $\rightarrow$  label = 1
- Generate fake image from  $G \rightarrow$  label = 0
- compute loss and update D

b. Train Generator:

- Generate fake images from random choice
- label as real
- compute loss and update G

6. Repeat until generator produces realistic color images

END

\*Observation:

- The generator gradually learned to produce realistic color images
- The discriminator improved in distinguishing real and fake images
- Generated images were visually similar to the training dataset after sufficient epochs
- Training demonstrated the adversarial learning process of GAN effectively

\*Result:

Successfully implemented a deep convolutional GAN to generate complex color images

Lab12.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all Connect Share

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import torchvision.utils as vutils
import matplotlib.pyplot as plt
import numpy as np

# 1. Hyperparameters
batch_size = 128
latent_dim = 100
epochs = 10
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# 2. Data loading (CIFAR-10 color images)
transform = transforms.Compose([
    transforms.Resize(64),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])
dataset = datasets.CIFAR10(root='./data', download=True, transform=transform)
dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True)

# 3. Define Generator
class Generator(nn.Module):
    def __init__(self):
        super(Generator, self).__init__()
        self.model = nn.Sequential(
            nn.ConvTranspose2d(100, 512, 4, 1, 0, bias=False),
            nn.BatchNorm2d(512), nn.ReLU(True),
            nn.ConvTranspose2d(512, 256, 4, 2, 1, bias=False),
            nn.BatchNorm2d(256), nn.ReLU(True),
            nn.ConvTranspose2d(256, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128), nn.ReLU(True),
            nn.ConvTranspose2d(128, 64, 4, 2, 1, bias=False),
            nn.BatchNorm2d(64), nn.ReLU(True),
            nn.ConvTranspose2d(64, 3, 4, 2, 1, bias=False),
            nn.Tanh()
        )
    def forward(self, x):
        return self.model(x)
```



Lab12.ipynb



File Edit View Insert Runtime Tools Help

Commands + Code ▾ + Text ▾ Run all ▾

```
[ ] # 4. Define Discriminator
class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()
        self.model = nn.Sequential(
            nn.Conv2d(3, 64, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(64, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128), nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(128, 256, 4, 2, 1, bias=False),
            nn.BatchNorm2d(256), nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(256, 512, 4, 2, 1, bias=False),
            nn.BatchNorm2d(512), nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(512, 1, 4, 1, 0, bias=False),
            nn.Sigmoid()
        )
    def forward(self, x):
        return self.model(x).view(-1, 1)

# 5. Initialize models
netG, netD = Generator().to(device), Discriminator().to(device)
criterion = nn.BCELoss()
optimizerG = optim.Adam(netG.parameters(), lr=0.0002, betas=(0.5, 0.999))
optimizerD = optim.Adam(netD.parameters(), lr=0.0002, betas=(0.5, 0.999))

# Lists to store losses
d_losses = []
g_losses = []

# 6. Training loop
for epoch in range(epochs):
    for i, (real_imgs, _) in enumerate(dataloader):
        real_imgs = real_imgs.to(device)
        batch_size = real_imgs.size(0)

        # Real labels = 1, Fake labels = 0
        real_labels = torch.ones(batch_size, 1, device=device)
        fake_labels = torch.zeros(batch_size, 1, device=device)

        # Train Discriminator
        z = torch.randn(batch_size, latent_dim, 1, 1, device=device)
        fake_imgs = netG(z)
        real_loss = criterion(netD(real_imgs), real_labels)
        fake_loss = criterion(netD(fake_imgs.detach()), fake_labels)
        d_loss = real_loss + fake_loss

        optimizerD.zero_grad()
```

```
# Store losses
d_losses.append(d_loss.item())
g_losses.append(g_loss.item())

print(f"Epoch [{epoch+1}/{epochs}] D Loss: {d_loss:.4f}, G Loss: {g_loss:.4f}")

# 7. Plot Epoch Graph (Loss vs Epoch)
plt.figure(figsize=(8,5))
plt.plot(range(1, epochs+1), d_losses, label='Discriminator Loss', marker='o')
plt.plot(range(1, epochs+1), g_losses, label='Generator Loss', marker='o')
plt.title("Epoch vs Loss (GAN Training Progress)")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()

# 8. Generate samples
z = torch.randn(64, latent_dim, 1, 1, device=device)
fake_imgs = netG(z)
vutils.save_image(fake_imgs, 'dcgan_generated.png', normalize=True)
plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Generated Images")
plt.imshow(np.transpose(vutils.make_grid(fake_imgs[:64], padding=2, normalize=True).cpu(), (1,2,0)))
plt.show()
```

100% |████████| 170M/170M [00:08<00:00, 19.3MB/s]  
Epoch [1/10] D Loss: 0.3572, G Loss: 3.6326  
Epoch [2/10] D Loss: 0.2364, G Loss: 5.8390  
Epoch [3/10] D Loss: 0.4392, G Loss: 3.4715  
Epoch [4/10] D Loss: 0.1285, G Loss: 4.1865  
Epoch [5/10] D Loss: 0.1988, G Loss: 3.7253  
Epoch [6/10] D Loss: 0.2754, G Loss: 4.0325  
Epoch [7/10] D Loss: 0.1110, G Loss: 8.6499  
Epoch [8/10] D Loss: 0.0016, G Loss: 8.4825  
Epoch [9/10] D Loss: 0.1210, G Loss: 4.8809  
Epoch [10/10] D Loss: 0.0389, G Loss: 6.7671

Epoch vs Loss (GAN Training Progress)

